

```
henry@hdp2-os-nn-01:~/myapp/hadoop-2.2.0> hadoop --config $HADOOP_HOME/etc/hadoop jar
./share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar -input in0/credit_card_tx_0_out.txt -output sout7
-mapper 'cat' -reducer 'wc -l' -verbose
STREAM: addTaskEnvironment=HADOOP_ROOT_LOGGER=
STREAM: shippedCanonFiles_=[]
STREAM: shipped: false /home/henry/myapp/hadoop-2.2.0/cat
STREAM: cmd=cat
STREAM: cmd=null
STREAM: shipped: false /home/henry/myapp/hadoop-2.2.0/wc
STREAM: cmd=wc -l
STREAM: Found runtime classes in: /tmp/hadoop-henry/hadoop-unjar7953759029010301842/
STREAM: ==== JobConf properties:
STREAM: dfs.block.access.key.update.interval=600
STREAM: dfs.block.access.token.enable=false
STREAM: dfs.block.access.token.lifetime=600
STREAM: dfs.blockreport.initialDelay=0
STREAM: dfs.blockreport.intervalMsec=21600000
STREAM: dfs.blocksize=134217728
STREAM: dfs.bytes-per-checksum=512
STREAM: dfs.client.write.packet.size=65536
STREAM: dfs.client.block.write.replace-datanode-on-failure.enable=true
STREAM: dfs.client.block.write.replace-datanode-on-failure.policy=DEFAULT
STREAM: dfs.client.block.write.retries=3
STREAM: dfs.client.failover.connection.retries=0
STREAM: dfs.client.failover.connection.retries.on.timeouts=0
STREAM: dfs.client.failover.max.attempts=15
STREAM: dfs.client.failover.sleep.base.millis=500
STREAM: dfs.client.failover.sleep.max.millis=15000
STREAM: dfs.client.file-block-storage-locations.num-threads=10
STREAM: dfs.client.file-block-storage-locations.timeout=60
STREAM: dfs.client.https.keystore.resource=ssl-client.xml
STREAM: dfs.client.https.need-auth=false
STREAM: dfs.client.use.datanode.hostname=false
STREAM: dfs.client.write.exclude.nodes.cache.expiry.interval.millis=600000
STREAM: dfs.datanode.address=0.0.0.0:50010
STREAM: dfs.datanode.available-space-volume-choosing-policy.balanced-space-preference-
fraction=0.75f
STREAM: dfs.datanode.available-space-volume-choosing-policy.balanced-space-
threshold=10737418240
STREAM: dfs.datanode.balance.bandwidthPerSec=1048576
STREAM: dfs.datanode.data.dir=file://${hadoop.tmp.dir}/dfs/data
STREAM: dfs.datanode.data.dir.perm=700
```

STREAM: dfs.datanode.directoryscan.interval=21600  
STREAM: dfs.datanode.directoryscan.threads=1  
STREAM: dfs.datanode.dns.interface=default  
STREAM: dfs.datanode.dns.nameserver=default  
STREAM: dfs.datanode.drop.cache.behind.reads=false  
STREAM: dfs.datanode.drop.cache.behind.writes=false  
STREAM: dfs.datanode.du.reserved=0  
STREAM: dfs.datanode.failed.volumes.tolerated=0  
STREAM: dfs.datanode.handler.count=10  
STREAM: dfs.datanode.hdfs-blocks-metadata.enabled=false  
STREAM: dfs.datanode.http.address=0.0.0.0:50075  
STREAM: dfs.datanode.https.address=0.0.0.0:50475  
STREAM: dfs.datanode.ipc.address=0.0.0.0:50020  
STREAM: dfs.datanode.max.transfer.threads=4096  
STREAM: dfs.datanode.readahead.bytes=4193404  
STREAM: dfs.datanode.sync.behind.writes=false  
STREAM: dfs.datanode.use.datanode.hostname=false  
STREAM: dfs.default.chunk.view.size=32768  
STREAM: dfs.encrypt.data.transfer=false  
STREAM: dfs.ha.automatic-failover.enabled=false  
STREAM: dfs.ha.fencing.ssh.connect-timeout=30000  
STREAM: dfs.ha.log-roll.period=120  
STREAM: dfs.ha.tail-edits.period=60  
STREAM: dfs.heartbeat.interval=3  
STREAM: dfs.https.enable=false  
STREAM: dfs.https.server.keystore.resource=ssl-server.xml  
STREAM: dfs.image.compress=false  
STREAM: dfs.image.compression.codec=org.apache.hadoop.io.compress.DefaultCodec  
STREAM: dfs.image.transfer.bandwidthPerSec=0  
STREAM: dfs.image.transfer.timeout=600000  
STREAM: dfs.journalnode.http-address=0.0.0.0:8480  
STREAM: dfs.journalnode.rpc-address=0.0.0.0:8485  
STREAM: dfs.namenode.access.time.precision=3600000  
STREAM: dfs.namenode.audit.loggers=default  
STREAM: dfs.namenode.avoid.read.stale.datanode=false  
STREAM: dfs.namenode.avoid.write.stale.datanode=false  
STREAM: dfs.namenode.backup.address=0.0.0.0:50100  
STREAM: dfs.namenode.backup.http-address=0.0.0.0:50105  
STREAM: dfs.namenode.checkpoint.check.period=60  
STREAM: dfs.namenode.checkpoint.dir=file://\${hadoop.tmp.dir}/dfs/namesecondary  
STREAM: dfs.namenode.checkpoint.edits.dir=\${dfs.namenode.checkpoint.dir}  
STREAM: dfs.namenode.checkpoint.max-retries=3

STREAM: dfs.namenode.checkpoint.period=3600  
STREAM: dfs.namenode.checkpoint.txns=1000000  
STREAM: dfs.namenode.decommission.interval=30  
STREAM: dfs.namenode.decommission.nodes.per.interval=5  
STREAM: dfs.namenode.delegation.key.update-interval=86400000  
STREAM: dfs.namenode.delegation.token.max-lifetime=604800000  
STREAM: dfs.namenode.delegation.token.renew-interval=86400000  
STREAM: dfs.namenode.edits.dir=\${dfs.namenode.name.dir}  
STREAM: dfs.namenode.edits.journal-  
plugin.qjournal=org.apache.hadoop.hdfs.qjournal.client.QuorumJournalManager  
STREAM: dfs.namenode.edits.noeditlogchannelflush=false  
STREAM: dfs.namenode.enable.retrycache=true  
STREAM: dfs.namenode.fs-limits.max-blocks-per-file=1048576  
STREAM: dfs.namenode.fs-limits.max-component-length=0  
STREAM: dfs.namenode.fs-limits.max-directory-items=0  
STREAM: dfs.namenode.fs-limits.min-block-size=1048576  
STREAM: dfs.namenode.handler.count=10  
STREAM: dfs.namenode.http-address=0.0.0.0:50070  
STREAM: dfs.namenode.https-address=0.0.0.0:50470  
STREAM: dfs.namenode.invalidate.work.pct.per.iteration=0.32f  
STREAM:  
dfs.namenode.kerberos.internal.spnego.principal=\${dfs.web.authentication.kerberos.principal}  
STREAM: dfs.namenode.logging.level=info  
STREAM: dfs.namenode.max.extra.edits.segments.retained=10000  
STREAM: dfs.namenode.max.objects=0  
STREAM: dfs.namenode.name.dir=file://\${hadoop.tmp.dir}/dfs/name  
STREAM: dfs.namenode.name.dir.restore=false  
STREAM: dfs.namenode.num.checkpoints.retained=2  
STREAM: dfs.namenode.num.extra.edits.retained=1000000  
STREAM: dfs.namenode.replication.considerLoad=true  
STREAM: dfs.namenode.replication.interval=3  
STREAM: dfs.namenode.replication.min=1  
STREAM: dfs.namenode.replication.work.multiplier.per.iteration=2  
STREAM: dfs.namenode.retrycache.expirytime.millis=600000  
STREAM: dfs.namenode.retrycache.heap.percent=0.03f  
STREAM: dfs.namenode.safemode.extension=30000  
STREAM: dfs.namenode.safemode.min.datanodes=0  
STREAM: dfs.namenode.safemode.threshold-pct=0.999f  
STREAM: dfs.namenode.secondary.http-address=0.0.0.0:50090  
STREAM: dfs.namenode.stale.datanode.interval=30000  
STREAM: dfs.namenode.support.allow.format=true  
STREAM: dfs.namenode.write.stale.datanode.ratio=0.5f

STREAM: dfs.permissions.enabled=true  
STREAM: dfs.permissions.superusergroup=supergroup  
STREAM: dfs.replication=3  
STREAM: dfs.replication.max=512  
STREAM:  
dfs.secondary.namenode.kerberos.internal.spnego.principal=\${dfs.web.authentication.kerberos.principal}  
STREAM: dfs.stream-buffer-size=4096  
STREAM: dfs.support.append=true  
STREAM: dfs.webhdfs.enabled=false  
STREAM: file.blocksize=67108864  
STREAM: file.bytes-per-checksum=512  
STREAM: file.client-write-packet-size=65536  
STREAM: file.replication=1  
STREAM: file.stream-buffer-size=4096  
STREAM: fs.AbstractFileSystem.file.impl=org.apache.hadoop.fs.local.LocalFs  
STREAM: fs.AbstractFileSystem.hdfs.impl=org.apache.hadoop.fs.Hdfs  
STREAM: fs.AbstractFileSystem.viewfs.impl=org.apache.hadoop.fs.viewfs.ViewFs  
STREAM: fs.automatic.close=true  
STREAM: fs.client.resolve.remote.symlinks=true  
STREAM: fs.defaultFS=file:///   
STREAM: fs.df.interval=60000  
STREAM: fs.ftp.host=0.0.0.0  
STREAM: fs.ftp.host.port=21  
STREAM: fs.permissions.umask-mode=022  
STREAM: fs.s3.block.size=67108864  
STREAM: fs.s3.buffer.dir=\${hadoop.tmp.dir}/s3  
STREAM: fs.s3.maxRetries=4  
STREAM: fs.s3.sleepTimeSeconds=10  
STREAM: fs.s3n.block.size=67108864  
STREAM: fs.trash.checkpoint.interval=0  
STREAM: fs.trash.interval=0  
STREAM: ftp.blocksize=67108864  
STREAM: ftp.bytes-per-checksum=512  
STREAM: ftp.client-write-packet-size=65536  
STREAM: ftp.replication=3  
STREAM: ftp.stream-buffer-size=4096  
STREAM: ha.failover-controller.cli-check.rpc-timeout.ms=20000  
STREAM: ha.failover-controller.graceful-fence.connection.retries=1  
STREAM: ha.failover-controller.graceful-fence.rpc-timeout.ms=5000  
STREAM: ha.failover-controller.new-active.rpc-timeout.ms=60000  
STREAM: ha.health-monitor.check-interval.ms=1000

STREAM: ha.health-monitor.connect-retry-interval.ms=1000  
STREAM: ha.health-monitor.rpc-timeout.ms=45000  
STREAM: ha.health-monitor.sleep-after-disconnect.ms=1000  
STREAM: ha.zookeeper.acl=world:anyone:rwcd  
STREAM: ha.zookeeper.parent-znode=/hadoop-ha  
STREAM: ha.zookeeper.session-timeout.ms=5000  
STREAM: hadoop.common.configuration.version=0.23.0  
STREAM: hadoop.fuse.connection.timeout=300  
STREAM: hadoop.fuse.timer.period=5  
STREAM: hadoop.hdfs.configuration.version=1  
STREAM: hadoop.http.authentication.kerberos.keytab=\${user.home}/hadoop.keytab  
STREAM: hadoop.http.authentication.kerberos.principal=HTTP/\_HOST@LOCALHOST  
STREAM: hadoop.http.authentication.signature.secret.file=\${user.home}/hadoop-http-auth-signature-secret  
STREAM: hadoop.http.authentication.simple.anonymous.allowed=true  
STREAM: hadoop.http.authentication.token.validity=36000  
STREAM: hadoop.http.authentication.type=simple  
STREAM: hadoop.http.filter.initializers=org.apache.hadoop.http.lib.StaticUserWebFilter  
STREAM: hadoop.http.staticuser.user=dr.who  
STREAM: hadoop.jetty.logs.serve.aliases=true  
STREAM: hadoop.kerberos.kinit.command=kinit  
STREAM: hadoop.rpc.protection=authentication  
STREAM: hadoop.rpc.socket.factory.class.default=org.apache.hadoop.net.StandardSocketFactory  
STREAM: hadoop.security.authentication=simple  
STREAM: hadoop.security.authorization=false  
STREAM:  
hadoop.security.group.mapping=org.apache.hadoop.security.JniBasedUnixGroupsMappingWithFallback  
STREAM: hadoop.security.group.mapping.ldap.directory.search.timeout=10000  
STREAM: hadoop.security.group.mapping.ldap.search.attr.group.name=cn  
STREAM: hadoop.security.group.mapping.ldap.search.attr.member=member  
STREAM: hadoop.security.group.mapping.ldap.search.filter.group=(objectClass=group)  
STREAM:  
hadoop.security.group.mapping.ldap.search.filter.user=(&(objectClass=user)(sAMAccountName={0}))  
STREAM: hadoop.security.group.mapping.ldap.ssl=false  
STREAM: hadoop.security.groups.cache.secs=300  
STREAM: hadoop.security.instrumentation.requires.admin=false  
STREAM: hadoop.security.uid.cache.secs=14400  
STREAM: hadoop.ssl.client.conf=ssl-client.xml  
STREAM: hadoop.ssl.enabled=false  
STREAM: hadoop.ssl.hostname.verifier=DEFAULT  
STREAM: hadoop.ssl.keystores.factory.class=org.apache.hadoop.security.ssl.FileBasedKeyStoresFactory  
STREAM: hadoop.ssl.require.client.cert=false

STREAM: hadoop.ssl.server.conf=ssl-server.xml  
STREAM: hadoop.tmp.dir=/tmp/hadoop-\${user.name}  
STREAM: hadoop.util.hash.type=murmur  
STREAM: hadoop.work.around.non.threadsafe.getpwuid=false  
STREAM: io.compression.codec.bzip2.library=system-native  
STREAM: io.file.buffer.size=4096  
STREAM: io.map.index.interval=128  
STREAM: io.map.index.skip=0  
STREAM: io.mapfile.bloom.error.rate=0.005  
STREAM: io.mapfile.bloom.size=1048576  
STREAM: io.native.lib.available=true  
STREAM: io.seqfile.compress.blocksize=1000000  
STREAM: io.seqfile.lazydecompress=true  
STREAM: io.seqfile.local.dir=\${hadoop.tmp.dir}/io/local  
STREAM: io.seqfile.sorter.recordlimit=1000000  
STREAM:  
io.serializations=org.apache.hadoop.io.serializer.WritableSerialization,org.apache.hadoop.io.serializer.avro.AvroSpecificSerialization,org.apache.hadoop.io.serializer.avro.AvroReflectSerialization  
STREAM: io.skip.checksum.errors=false  
STREAM: ipc.client.connect.max.retries=10  
STREAM: ipc.client.connect.max.retries.on.timeouts=45  
STREAM: ipc.client.connect.timeout=20000  
STREAM: ipc.client.connection.maxidletime=10000  
STREAM: ipc.client.fallback-to-simple-auth-allowed=false  
STREAM: ipc.client.idlethreshold=4000  
STREAM: ipc.client.kill.max=10  
STREAM: ipc.client.tcpnodelay=false  
STREAM: ipc.server.listen.queue.size=128  
STREAM: ipc.server.tcpnodelay=false  
STREAM: map.sort.class=org.apache.hadoop.util.QuickSort  
STREAM: mapred.child.java.opts=-Xmx200m  
STREAM: mapred.input.dir=file:/home/henry/myapp/hadoop-2.2.0/in0/credit\_card\_tx\_0\_out.txt  
STREAM: mapred.input.format.class=org.apache.hadoop.mapred.TextInputFormat  
STREAM: mapred.jar=/home/henry/myapp/hadoop-2.2.0./share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar  
STREAM: mapred.map.runner.class=org.apache.hadoop.streaming.PipeMapRunner  
STREAM: mapred.mapoutput.key.class=org.apache.hadoop.io.Text  
STREAM: mapred.mapoutput.value.class=org.apache.hadoop.io.Text  
STREAM: mapred.mapper.class=org.apache.hadoop.streaming.PipeMapper  
STREAM: mapred.output.dir=file:/home/henry/myapp/hadoop-2.2.0/out7  
STREAM: mapred.output.format.class=org.apache.hadoop.mapred.TextOutputFormat  
STREAM: mapred.output.key.class=org.apache.hadoop.io.Text

STREAM: mapred.output.value.class=org.apache.hadoop.io.Text  
STREAM: mapred.reducer.class=org.apache.hadoop.streaming.PipeReducer  
STREAM: mapred.working.dir=file:/home/henry/myapp/hadoop-2.2.0  
STREAM: mapreduce.admin.user.env=LD\_LIBRARY\_PATH=\$HADOOP\_COMMON\_HOME/lib/native  
STREAM: mapreduce.am.max-attempts=2  
STREAM: mapreduce.client.completion.pollinterval=5000  
STREAM: mapreduce.client.genericoptionsparser.used=true  
STREAM: mapreduce.client.output.filter=FAILED  
STREAM: mapreduce.client.progressmonitor.pollinterval=1000  
STREAM: mapreduce.client.submit.file.replication=10  
STREAM: mapreduce.cluster.acls.enabled=false  
STREAM: mapreduce.cluster.local.dir=\${hadoop.tmp.dir}/mapred/local  
STREAM: mapreduce.cluster.temp.dir=\${hadoop.tmp.dir}/mapred/temp  
STREAM: mapreduce.framework.name=local  
STREAM: mapreduce.ifile.readahead=true  
STREAM: mapreduce.ifile.readahead.bytes=4194304  
STREAM: mapreduce.input.fileinputformat.inputdir=file:/home/henry/myapp/hadoop-2.2.0/in0/credit\_card\_tx\_0\_out.txt  
STREAM: mapreduce.input.fileinputformat.split.minsize=0  
STREAM: mapreduce.job.acl-modify-job=  
STREAM: mapreduce.job.acl-view-job=  
STREAM: mapreduce.job.classloader=false  
STREAM:  
mapreduce.job.classloader.system.classes=java.,javax.,org.apache.commons.logging.,org.apache.log4j.,org.apache.hadoop.  
STREAM: mapreduce.job.committer.setup.cleanup.needed=true  
STREAM: mapreduce.job.complete.cancel.delegation.tokens=true  
STREAM: mapreduce.job.counters.max=120  
STREAM: mapreduce.job.end-notification.max.attempts=5  
STREAM: mapreduce.job.end-notification.max.retry.interval=5000  
STREAM: mapreduce.job.end-notification.retry.attempts=0  
STREAM: mapreduce.job.end-notification.retry.interval=1000  
STREAM: mapreduce.job.hdfs-servers=\${fs.defaultFS}  
STREAM: mapreduce.job.jar=/home/henry/myapp/hadoop-2.2.0/./share/hadoop/tools/lib/hadoop-streaming-2.2.0.jar  
STREAM: mapreduce.job.jvm.numtasks=1  
STREAM:  
mapreduce.job.map.output.collector.class=org.apache.hadoop.mapred.MapTask\$MapOutputBuffer  
STREAM: mapreduce.job.maps=2  
STREAM: mapreduce.job.maxtaskfailures.per.tracker=3  
STREAM: mapreduce.job.output.key.class=org.apache.hadoop.io.Text  
STREAM: mapreduce.job.output.value.class=org.apache.hadoop.io.Text

STREAM: mapreduce.job.queueName=default  
STREAM:  
mapreduce.job.reduce.shuffle.consumer.plugin.class=org.apache.hadoop.mapreduce.task.reduce.Shuffle  
e  
STREAM: mapreduce.job.reduce.slowstart.completedmaps=0.05  
STREAM: mapreduce.job.reduces=1  
STREAM: mapreduce.job.speculative.slowNodeThreshold=1.0  
STREAM: mapreduce.job.speculative.slowTaskThreshold=1.0  
STREAM: mapreduce.job.speculative.speculativeCap=0.1  
STREAM: mapreduce.job.split.metainfo.maxSize=10000000  
STREAM: mapreduce.job.token.tracking.ids.enabled=false  
STREAM: mapreduce.job.ubertask.enable=false  
STREAM: mapreduce.job.ubertask.maxMaps=9  
STREAM: mapreduce.job.ubertask.maxReduces=1  
STREAM: mapreduce.job.userlog.retain.hours=24  
STREAM: mapreduce.job.working.dir=file:/home/henry/myapp/hadoop-2.2.0  
STREAM: mapreduce.jobhistory.address=0.0.0.0:10020  
STREAM: mapreduce.jobhistory.cleaner.enable=true  
STREAM: mapreduce.jobhistory.cleaner.interval-ms=86400000  
STREAM: mapreduce.jobhistory.client.thread-count=10  
STREAM: mapreduce.jobhistory.datestring.cache.size=200000  
STREAM: mapreduce.jobhistory.done-dir=\${yarn.app.mapreduce.am.staging-dir}/history/done  
STREAM: mapreduce.jobhistory.http.policy=HTTP\_ONLY  
STREAM: mapreduce.jobhistory.intermediate-done-dir=\${yarn.app.mapreduce.am.staging-dir}/history/done\_intermediate  
STREAM: mapreduce.jobhistory.joblist.cache.size=20000  
STREAM: mapreduce.jobhistory.keytab=/etc/security/keytab/jhs.service.keytab  
STREAM: mapreduce.jobhistory.loadedJobs.cache.size=5  
STREAM: mapreduce.jobhistory.max-age-ms=604800000  
STREAM: mapreduce.jobhistory.miniCluster.fixed.ports=false  
STREAM: mapreduce.jobhistory.move.interval-ms=180000  
STREAM: mapreduce.jobhistory.move.thread-count=3  
STREAM: mapreduce.jobhistory.principal=jhs/\_HOST@REALM.TLD  
STREAM: mapreduce.jobhistory.webapp.address=0.0.0.0:19888  
STREAM: mapreduce.jobtracker.address=local  
STREAM: mapreduce.jobtracker.expire.trackers.interval=600000  
STREAM: mapreduce.jobtracker.handler.count=10  
STREAM: mapreduce.jobtracker.heartbeats.in.second=100  
STREAM: mapreduce.jobtracker.http.address=0.0.0.0:50030  
STREAM: mapreduce.jobtracker.instrumentation=org.apache.hadoop.mapred.JobTrackerMetricsInst  
STREAM: mapreduce.jobtracker.jobhistory.block.size=3145728  
STREAM: mapreduce.jobtracker.jobhistory.lru.cache.size=5



STREAM: mapreduce.jobtracker.jobhistory.task.numberprogresssplits=12  
STREAM: mapreduce.jobtracker.maxtasks.perjob=-1  
STREAM: mapreduce.jobtracker.persist.jobstatus.active=true  
STREAM: mapreduce.jobtracker.persist.jobstatus.dir=/jobtracker/jobsInfo  
STREAM: mapreduce.jobtracker.persist.jobstatus.hours=1  
STREAM: mapreduce.jobtracker.restart.recover=false  
STREAM: mapreduce.jobtracker.retiredjobs.cache.size=1000  
STREAM: mapreduce.jobtracker.staging.root.dir=\${hadoop.tmp.dir}/mapred/staging  
STREAM: mapreduce.jobtracker.system.dir=\${hadoop.tmp.dir}/mapred/system  
STREAM: mapreduce.jobtracker.taskcache.levels=2  
STREAM: mapreduce.jobtracker.taskscheduler=org.apache.hadoop.mapred.JobQueueTaskScheduler  
STREAM: mapreduce.jobtracker.tasktracker.maxblacklists=4  
STREAM: mapreduce.local.clientfactory.class.name=org.apache.hadoop.mapred.LocalClientFactory  
STREAM: mapreduce.map.cpu.vcores=1  
STREAM: mapreduce.map.log.level=INFO  
STREAM: mapreduce.map.maxattempts=4  
STREAM: mapreduce.map.output.compress=false  
STREAM: mapreduce.map.output.compress.codec=org.apache.hadoop.io.compress.DefaultCodec  
STREAM: mapreduce.map.output.key.class=org.apache.hadoop.io.Text  
STREAM: mapreduce.map.output.value.class=org.apache.hadoop.io.Text  
STREAM: mapreduce.map.skip.maxrecords=0  
STREAM: mapreduce.map.skip.proc.count.autoincr=true  
STREAM: mapreduce.map.sort.spill.percent=0.80  
STREAM: mapreduce.map.speculative=true  
STREAM: mapreduce.output.fileoutputformat.compress=false  
STREAM:  
mapreduce.output.fileoutputformat.compress.codec=org.apache.hadoop.io.compress.DefaultCodec  
STREAM: mapreduce.output.fileoutputformat.compress.type=RECORD  
STREAM: mapreduce.output.fileoutputformat.outputdir=file:/home/henry/myapp/hadoop-2.2.0/sout7  
STREAM: mapreduce.reduce.cpu.vcores=1  
STREAM: mapreduce.reduce.input.buffer.percent=0.0  
STREAM: mapreduce.reduce.log.level=INFO  
STREAM: mapreduce.reduce.markreset.buffer.percent=0.0  
STREAM: mapreduce.reduce.maxattempts=4  
STREAM: mapreduce.reduce.merge.inmem.threshold=1000  
STREAM: mapreduce.reduce.shuffle.connect.timeout=180000  
STREAM: mapreduce.reduce.shuffle.input.buffer.percent=0.70  
STREAM: mapreduce.reduce.shuffle.memory.limit.percent=0.25  
STREAM: mapreduce.reduce.shuffle.merge.percent=0.66  
STREAM: mapreduce.reduce.shuffle.parallelcopies=5  
STREAM: mapreduce.reduce.shuffle.read.timeout=180000  
STREAM: mapreduce.reduce.shuffle.retry-delay.max.ms=60000

STREAM: mapreduce.reduce.skip.maxgroups=0  
STREAM: mapreduce.reduce.skip.proc.count.autoincr=true  
STREAM: mapreduce.reduce.speculative=true  
STREAM: mapreduce.shuffle.max.connections=0  
STREAM: mapreduce.shuffle.port=13562  
STREAM: mapreduce.shuffle.ssl.enabled=false  
STREAM: mapreduce.shuffle.ssl.file.buffer.size=65536  
STREAM: mapreduce.task.files.preserve.failedtasks=false  
STREAM: mapreduce.task.io.sort.factor=10  
STREAM: mapreduce.task.io.sort.mb=100  
STREAM: mapreduce.task.merge.progress.records=10000  
STREAM: mapreduce.task.profile=false  
STREAM: mapreduce.task.profile.maps=0-2  
STREAM: mapreduce.task.profile.reduces=0-2  
STREAM: mapreduce.task.skip.start.attempts=2  
STREAM: mapreduce.task.timeout=600000  
STREAM: mapreduce.task.tmp.dir=./tmp  
STREAM: mapreduce.task.userlog.limit.kb=0  
STREAM: mapreduce.tasktracker.dns.interface=default  
STREAM: mapreduce.tasktracker.dns.nameserver=default  
STREAM: mapreduce.tasktracker.healthchecker.interval=60000  
STREAM: mapreduce.tasktracker.healthchecker.script.timeout=600000  
STREAM: mapreduce.tasktracker.http.address=0.0.0.0:50060  
STREAM: mapreduce.tasktracker.http.threads=40  
STREAM: mapreduce.tasktracker.indexcache.mb=10  
STREAM: mapreduce.tasktracker.instrumentation=org.apache.hadoop.mapred.TaskTrackerMetricsInst  
STREAM: mapreduce.tasktracker.local.dir.minspacekill=0  
STREAM: mapreduce.tasktracker.local.dir.minspacestart=0  
STREAM: mapreduce.tasktracker.map.tasks.maximum=2  
STREAM: mapreduce.tasktracker.outofband.heartbeat=false  
STREAM: mapreduce.tasktracker.reduce.tasks.maximum=2  
STREAM: mapreduce.tasktracker.report.address=127.0.0.1:0  
STREAM: mapreduce.tasktracker.taskcontroller=org.apache.hadoop.mapred.DefaultTaskController  
STREAM: mapreduce.tasktracker.taskmemorymanager.monitoringinterval=5000  
STREAM: mapreduce.tasktracker.tasks.sleeptimebeforekill=5000  
STREAM: net.topology.impl=org.apache.hadoop.net.NetworkTopology  
STREAM: net.topology.node.switch.mapping.impl=org.apache.hadoop.net.ScriptBasedMapping  
STREAM: net.topology.script.number.args=100  
STREAM: nfs3.mountd.port=4242  
STREAM: nfs3.server.port=2049  
STREAM: s3.blocksize=67108864  
STREAM: s3.bytes-per-checksum=512

STREAM: s3.client-write-packet-size=65536  
STREAM: s3.replication=3  
STREAM: s3.stream-buffer-size=4096  
STREAM: s3native.blocksize=67108864  
STREAM: s3native.bytes-per-checksum=512  
STREAM: s3native.client-write-packet-size=65536  
STREAM: s3native.replication=3  
STREAM: s3native.stream-buffer-size=4096  
STREAM: stream.addenvironment=HADOOP\_ROOT\_LOGGER=  
STREAM: stream.map.input.writer.class=org.apache.hadoop.streaming.io.TextInputWriter  
STREAM: stream.map.output.reader.class=org.apache.hadoop.streaming.io.TextOutputReader  
STREAM: stream.map.streamprocessor=cat  
STREAM: stream.reduce.input.writer.class=org.apache.hadoop.streaming.io.TextInputWriter  
STREAM: stream.reduce.output.reader.class=org.apache.hadoop.streaming.io.TextOutputReader  
STREAM: stream.reduce.streamprocessor=wc+-l  
STREAM: tfile.fs.input.buffer.size=262144  
STREAM: tfile.fs.output.buffer.size=262144  
STREAM: tfile.io.chunk.size=1048576  
STREAM: yarn.acl.enable=true  
STREAM: yarn.admin.acl=  
STREAM: yarn.am.liveness-monitor.expiry-interval-ms=600000  
STREAM: yarn.app.mapreduce.am.command-opts=-Xmx1024m  
STREAM: yarn.app.mapreduce.am.job.committer.cancel-timeout=60000  
STREAM: yarn.app.mapreduce.am.job.committer.commit-window=10000  
STREAM: yarn.app.mapreduce.am.job.task.listener.thread-count=30  
STREAM: yarn.app.mapreduce.am.resource.cpu-vcores=1  
STREAM: yarn.app.mapreduce.am.resource.mb=1536  
STREAM: yarn.app.mapreduce.am.scheduler.heartbeat.interval-ms=1000  
STREAM: yarn.app.mapreduce.am.staging-dir=/tmp/hadoop-yarn/staging  
STREAM: yarn.app.mapreduce.client-am.ipc.max-retries=3  
STREAM: yarn.app.mapreduce.client.max-retries=3  
STREAM:  
yarn.application.classpath=\$HADOOP\_CONF\_DIR,\$HADOOP\_COMMON\_HOME/share/hadoop/common/\*,\$HADOOP\_COMMON\_HOME/share/hadoop/common/lib/\*,\$HADOOP\_HDFS\_HOME/share/hadoop/hdfs/\*,\$HADOOP\_HDFS\_HOME/share/hadoop/hdfs/lib/\*,\$HADOOP\_YARN\_HOME/share/hadoop/yarn/\*,\$HADOOP\_YARN\_HOME/share/hadoop/yarn/lib/\*  
STREAM: yarn.client.app-submission.poll-interval=1000  
STREAM: yarn.client.max-nodemanagers-proxies=500  
STREAM: yarn.client.nodemanager-client-async.thread-pool-max-size=500  
STREAM: yarn.http.policy=HTTP\_ONLY  
STREAM: yarn.ipc.rpc.class=org.apache.hadoop.yarn.ipc.HadoopYarnProtoRPC  
STREAM: yarn.ipc.serializer.type=protocolbuffers

STREAM: yarn.log-aggregation-enable=false  
STREAM: yarn.log-aggregation.retain-check-interval-seconds=-1  
STREAM: yarn.log-aggregation.retain-seconds=-1  
STREAM: yarn.nm.liveness-monitor.expiry-interval-ms=600000  
STREAM: yarn.nodemanager.address=\${yarn.nodemanager.hostname}:0  
STREAM: yarn.nodemanager.admin-env=MALLOC\_ARENA\_MAX=\$MALLOC\_ARENA\_MAX  
STREAM: yarn.nodemanager.aux-  
services.mapreduce\_shuffle.class=org.apache.hadoop.mapred.ShuffleHandler  
STREAM: yarn.nodemanager.container-  
executor.class=org.apache.hadoop.yarn.server.nodemanager.DefaultContainerExecutor  
STREAM: yarn.nodemanager.container-manager.thread-count=20  
STREAM: yarn.nodemanager.container-monitor.interval-ms=3000  
STREAM: yarn.nodemanager.delete.debug-delay-sec=0  
STREAM: yarn.nodemanager.delete.thread-count=4  
STREAM: yarn.nodemanager.disk-health-checker.interval-ms=120000  
STREAM: yarn.nodemanager.disk-health-checker.min-healthy-disks=0.25  
STREAM: yarn.nodemanager.env-  
whitelist=JAVA\_HOME,HADOOP\_COMMON\_HOME,HADOOP\_HDFS\_HOME,HADOOP\_CONF\_DIR,HADO  
OP\_YARN\_HOME  
STREAM: yarn.nodemanager.health-checker.interval-ms=600000  
STREAM: yarn.nodemanager.health-checker.script.timeout-ms=1200000  
STREAM: yarn.nodemanager.hostname=0.0.0.0  
STREAM: yarn.nodemanager.keytab=/etc/krb5.keytab  
STREAM: yarn.nodemanager.linux-container-executor.cgroups.hierarchy=/hadoop-yarn  
STREAM: yarn.nodemanager.linux-container-executor.cgroups.mount=false  
STREAM: yarn.nodemanager.linux-container-executor.resources-  
handler.class=org.apache.hadoop.yarn.server.nodemanager.util.DefaultLCEResourcesHandler  
STREAM: yarn.nodemanager.local-cache.max-files-per-directory=8192  
STREAM: yarn.nodemanager.local-dirs=\${hadoop.tmp.dir}/nm-local-dir  
STREAM: yarn.nodemanager.localizer.address=\${yarn.nodemanager.hostname}:8040  
STREAM: yarn.nodemanager.localizer.cache.cleanup.interval-ms=600000  
STREAM: yarn.nodemanager.localizer.cache.target-size-mb=10240  
STREAM: yarn.nodemanager.localizer.client.thread-count=5  
STREAM: yarn.nodemanager.localizer.fetch.thread-count=4  
STREAM: yarn.nodemanager.log-aggregation.compression-type=none  
STREAM: yarn.nodemanager.log-dirs=\${yarn.log.dir}/userlogs  
STREAM: yarn.nodemanager.log.retain-seconds=10800  
STREAM: yarn.nodemanager.pmem-check-enabled=true  
STREAM: yarn.nodemanager.process-kill-wait.ms=2000  
STREAM: yarn.nodemanager.remote-app-log-dir=/tmp/logs  
STREAM: yarn.nodemanager.remote-app-log-dir-suffix=logs  
STREAM: yarn.nodemanager.resource.cpu-vcores=8

STREAM: yarn.nodemanager.resource.memory-mb=8192  
STREAM: yarn.nodemanager.resourcemanager.connect.retry\_interval.secs=30  
STREAM: yarn.nodemanager.resourcemanager.connect.wait.secs=900  
STREAM: yarn.nodemanager.sleep-delay-before-sigkill.ms=250  
STREAM: yarn.nodemanager.vmem-check-enabled=true  
STREAM: yarn.nodemanager.vmem-pmem-ratio=2.1  
STREAM: yarn.nodemanager.webapp.address=\${yarn.nodemanager.hostname}:8042  
STREAM: yarn.resourcemanager.address=\${yarn.resourcemanager.hostname}:8032  
STREAM: yarn.resourcemanager.admin.address=\${yarn.resourcemanager.hostname}:8033  
STREAM: yarn.resourcemanager.admin.client.thread-count=1  
STREAM: yarn.resourcemanager.am.max-attempts=2  
STREAM: yarn.resourcemanager.amliveliness-monitor.interval-ms=1000  
STREAM: yarn.resourcemanager.application-tokens.master-key-rolling-interval-secs=86400  
STREAM: yarn.resourcemanager.client.thread-count=50  
STREAM: yarn.resourcemanager.connect.max-wait.ms=900000  
STREAM: yarn.resourcemanager.connect.retry-interval.ms=30000  
STREAM: yarn.resourcemanager.container-tokens.master-key-rolling-interval-secs=86400  
STREAM: yarn.resourcemanager.container.liveness-monitor.interval-ms=600000  
STREAM: yarn.resourcemanager.delayed.delegation-token.removal-interval-ms=30000  
STREAM: yarn.resourcemanager.fs.state-store.uri=\${hadoop.tmp.dir}/yarn/system/rmstore  
STREAM: yarn.resourcemanager.hostname=0.0.0.0  
STREAM: yarn.resourcemanager.keytab=/etc/krb5.keytab  
STREAM: yarn.resourcemanager.max-completed-applications=10000  
STREAM: yarn.resourcemanager.nm.liveness-monitor.interval-ms=1000  
STREAM: yarn.resourcemanager.nodemanagers.heartbeat-interval-ms=1000  
STREAM: yarn.resourcemanager.recovery.enabled=false  
STREAM: yarn.resourcemanager.resource-tracker.address=\${yarn.resourcemanager.hostname}:8031  
STREAM: yarn.resourcemanager.resource-tracker.client.thread-count=50  
STREAM: yarn.resourcemanager.scheduler.address=\${yarn.resourcemanager.hostname}:8030  
STREAM:  
yarn.resourcemanager.scheduler.class=org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.CapacityScheduler  
STREAM: yarn.resourcemanager.scheduler.client.thread-count=50  
STREAM: yarn.resourcemanager.scheduler.monitor.enable=false  
STREAM:  
yarn.resourcemanager.scheduler.monitor.policies=org.apache.hadoop.yarn.server.resourcemanager.monitor.capacity.ProportionalCapacityPreemptionPolicy  
STREAM:  
yarn.resourcemanager.store.class=org.apache.hadoop.yarn.server.resourcemanager.recovery.FileSystemRMStateStore  
STREAM: yarn.resourcemanager.webapp.address=\${yarn.resourcemanager.hostname}:8088  
STREAM: yarn.resourcemanager.webapp.https.address=\${yarn.resourcemanager.hostname}:8090

```
STREAM: yarn.scheduler.maximum-allocation-mb=8192  
STREAM: yarn.scheduler.maximum-allocation-vcores=32  
STREAM: yarn.scheduler.minimum-allocation-mb=1024  
STREAM: yarn.scheduler.minimum-allocation-vcores=1  
STREAM: ====
```

```
.....
```

```
henry@hdp2-os-nn-01:~/myapp/hadoop-2.2.0>
```